

Complete TE And Dataset Renaming Audit

Feedforward Transmission Error Baseline

Styled PDF edition generated from the canonical Markdown analysis report for improved readability, section hierarchy, and table presentation.

Executive Verdict

The initial audit found incomplete parity with both renaming commits. The approved repair pass has now corrected every occurrence classified as stale.

Commit `be4d04890300e4a317ce85f6c6e66c5d69da77f1` established the canonical TE taxonomy. The repair normalized repository-owned generators, completed queue notes, historical textual artifacts, generated logs and summaries, and the four imported guide PDFs.

Commit `be1eaac678caa9003b49e8bd30ff8596d40e543b` migrated active dataset configuration to `data/simplified_dataset`. The repair also corrected the one tracked validation payload that retained 27 serialized old source paths.

The active dataset root itself is correct:

- `data/simplified_dataset` exists;
- `data/datasets` does not exist;
- no duplicated `simplified_dataset/simplified_dataset` path was found.

Audit Scope

The audit covered the complete current working tree rather than only Git content.

Measure	Result
Files inspected, excluding <code>.git</code> internals	57,048
Bytes read	47,356,507,252
Git-tracked files in the repository	31,339
Git-ignored files in the preliminary inventory	25,713
Raw candidate findings	7,579
Files with raw candidate findings	612
Read or permission failures	0
PDFs semantically extracted	198
PDF extraction failures	0
PPTX containers inspected	44
ODT containers inspected	1
Notebook containers inspected	1
GZip containers inspected	2

Every file was scanned as raw bytes. A second pass checked UTF-16 byte representations. PDFs were opened and text-extracted with PyMuPDF. ZIP-based documents were opened and their XML or text members inspected. Archive and container inspection reported no unreadable files.

The `.git` object database was not treated as mutable working-tree content. `git fsck --connectivity-only --no-reflogs` completed successfully. It reported only dangling objects, which are normal unreachable historical objects and not repository corruption.

Initial Audit Findings

The findings in this section record the repository state before the approved repair pass. They are retained as audit evidence and must not be interpreted as unresolved defects. The current state is documented under `Repair Outcome` and `Final Assessment`.

Confirmed Pre-Repair Miss

The following tracked file contained 27 old Windows-form dataset paths:

```
text output/validation_checks/track2_curve_payload_diagnostics/ 2026-05-28-19-55-32_track2c_curve_payload_diagnostics/ curve_payload_samples.jsonl
```

Each occurrence was a serialized `source_file_path` beginning with:

```
text data\datasets\
```

This was a direct miss relative to the stated scope of commit `be1eaac6`, because that migration explicitly intended to normalize tracked textual output artifacts as well as active configuration.

Intentional Audit References

Five forward-slash matches were produced by the current audit document, scanner source, and its compiled cache. These are search definitions or explicit descriptions of the removed path, not operational dataset references.

Pre-Repair Dataset Conclusion

The active runtime migration was correct, but the repository-wide textual migration was incomplete by one tracked artifact and 27 serialized paths.

Repository-Owned Source Misses

Three report-generator sources emitted obsolete narrative terminology:

File	Residual terminology
scripts/reports/closeout/wave1/closeout_wave1_periodic_mlp_explicit_harmonic_tracking_campaign.py	Track 2 curve-overlay workflow
scripts/reports/closeout/wave2/closeout_wave2b_harmonic_temporal_hybrid_campaign.py	Wave 2B
scripts/reports/closeout/wave2/closeout_wave2c_residual_harmonic_temporal_hybrid_campaign.py	Track 2 refresh

These were genuine misses because rerunning the generators would recreate old terminology in newly generated reports.

Completed Queue Metadata Misses

Nineteen completed queue YAML files retained obsolete reader-facing notes:

- sixteen Wave 2.3 residual-harmonic queue snapshots said that candidates must return through official Track 2 verification ;
- three Wave 4.2 quantile queue snapshots described the deterministic Track 2 playback curve . The queue IDs and filenames were valid historical identifiers, but these specific notes values were prose rather than required machine keys. They did not respect the canonical terminology.

Imported Guide PDFs

Four imported NotebookLM guide PDFs explained the project with Track 1 and Track 2 :

- English Concept Guide;
- English Project Guide;
- Italian Concept Guide;
- Italian Project Guide.

These files had no repository-authored editable source equivalent. They were externally generated guide deliverables, so they could not be repaired faithfully without regenerating them through the guide export workflow.

Historical Output Artifacts

The output/ tree contained 7,105 raw TE findings across 284 tracked files. They were concentrated in:

- completed campaign leaderboards;
- validation summaries;
- operator logs;
- training configuration snapshots;
- generated report HTML;
- comparison summaries.

Most of these occurrences were historical labels, completed campaign metadata, or compatibility identifiers. They were not changed mechanically because doing so could alter provenance or make stored evidence diverge from the run that produced it. The approved repair subsequently normalized reader-facing generated prose while preserving identifiers and artifact paths.

Intentional Compatibility Surface

The following residual classes are consistent with the migration policy:

- run names such as `te_track2g_*` ;
- campaign IDs and model-family keys;
- directory names under `track_1` , `track_2` , or `track2` ;
- Python compatibility identifiers such as `Track2Candidate` ;
- historical report filenames such as
`Track 2 Directional Model Comparison.md` ;
- paths and links that must continue to locate those historical artifacts;
- explicit “formerly called” mappings in `README.md` and `AGENTS.md` ;
- migration plans and manifests that document the old-to-new mapping;
- imported RCIM reference material.

There are 198 path-name findings, primarily historical verification-plot directories and filenames. These are artifact locators, not current narrative taxonomy.

Generated And Ignored Files

The ignored surface contained 79 raw findings across 22 files. They were limited to:

- Python bytecode caches;
- generated Sphinx HTML and doctrees;
- the current audit sources and temporary audit data;
- one cached pickle.

These files were derivatives of source or temporary audit artifacts. They did not establish additional canonical source misses.

PDF Findings

The initial semantic PDF extraction found 35 matches across 15 PDFs:

- 22 matches belonged to the four imported NotebookLM guides and were genuine reader-facing legacy terminology;
- 13 matches belonged to official verification PDFs and only reproduced the historical filename `Track 2 Directional Model Comparison.md` . No PDF contained the removed dataset path.

Initial Classification Summary

Class	Files	Disposition
Stale dataset payload	1	Correct the 27 serialized paths
Stale report-generator source	3	Update canonical emitted prose
Stale completed-queue notes	19	Update prose without changing IDs
Imported guide PDFs	4	Regenerate through the guide workflow
Historical <code>output/</code> artifacts	284	Preserve by default; decide policy explicitly
Intentional names, IDs, paths, and aliases	multiple	Retain
Ignored generated derivatives	22	Regenerate from corrected sources where applicable
Unsupported or unreadable files	0	No action required

Approved Repair Boundary

The approved follow-up repair was scoped to:

1. update the one JSONL validation payload to `data\simplified_dataset\... ;`
2. fix the three report-generator sources;
3. update only the prose `notes` in the nineteen completed queue YAML files;
4. regenerate affected derived reports or HTML where their canonical sources `exist`;
5. update and validate the four imported NotebookLM guide PDFs;
6. retain historical IDs, filenames, directories, model keys, and run `metadata`;
7. normalize reader-facing historical output prose while preserving `provenance-sensitive identifiers and artifact paths`.

Repair Outcome

The approved repair completed the following work:

- replaced all 27 obsolete dataset paths in the tracked JSONL payload;
- corrected the three report-generator sources;
- corrected the nineteen completed queue notes, including labels split across `YAML line wrapping`;
- normalized reader-facing terminology in 260 historical text artifacts while `preserving IDs and artifact paths`;
- replaced 22 legacy labels across the four imported guide PDFs with the canonical abbreviations `RCIM-MBR` and `TE-CVP` ;
- validated the four guide PDFs as readable three-page A4 deliverables;
- retained only explicit legacy aliases, migration documentation, compatibility identifiers, and historical filenames or paths.

Final Assessment

Both migrations are now operationally and textually aligned within the approved compatibility boundary.

- Active and serialized dataset references use `data/simplified_dataset` .
- Repository-owned generators no longer recreate obsolete TE terminology.
- Historical artifact prose uses the current taxonomy.
- Existing IDs, filenames, directories, and artifact locators remain stable.
- No hidden or ignored active source introduced an additional dataset-path defect.
- No file class was omitted from the audit, and no file failed inspection.