

Wave 5.2R Stage 3 Analytical Anchor Reproduction And Stress Tests

Feedforward Transmission Error Baseline

Styled PDF edition generated from the canonical Markdown analysis report for improved readability, section hierarchy, and table presentation.

Executive Decision

Wave 5.2R Stage 3 passes all twelve exit gates.

PF_A_LOCAL_QUADRATIC is qualified as the analytical anchor for the next forward residual-learning stage, with one important boundary:

PF-A is trusted as a bounded interpolation component only inside the declared supported-core envelope.

The independent refit reproduces the Phase 1 model exactly:

- split signature: `c1aa8718fb9bf88cc2021c121dc4f3b4010fc1d2e45ac90af5f4376aa64f8e16` ;
- training curves used for fitting: 675 ;
- validation curves used only for evaluation: 194 ;
- test curves used only for evaluation: 97 ;
- design condition number: 5.361553 ;
- test raw MAE: 0.001807 deg ;
- test centered-shape MAE: 0.001385 deg ;
- test offset absolute error: 0.000965 deg ;
- maximum Phase 1 reproduction difference: 0.0 .

All 966 eligible forward conditions produce finite predictions. Sixty-four deterministic bootstrap refits remain stable, with a condition-number P95 of 5.612715 , a relative coefficient-change P95 of 0.029912 , and a prediction-deviation-to-anchor-MAE P95 of 0.196634 . The stress tests also expose limitations that must remain visible:

- withholding the lowest temperature level produces the worst axis holdout, with 0.004742 deg MAE and a 2.67 error ratio relative to the full-fit anchor on the same population;
- omitting the low-order harmonic group is the most damaging corruption, creating 0.010982 deg mean deviation from the anchor;
- validation and test each contain one condition outside at least one training-axis bound;
- reduced, PLC-safe, paper-order, and recovered ONNX formulations remain comparators only. No neural network training was executed.

Scope And Frozen Contract

Stage 3 is restricted to

- dataset: `polished_dataset` ;
- operating inputs: setpoint speed, setpoint torque, and temperature;
- direction: `Fw` ;
- analytical family: Polynomial-Fourier coefficient surfaces;
- fit partition: frozen training split only;
- evaluation partitions: unchanged validation and test splits;
- angular response: full-resolution transmission-error curve.

The analysis does not:

- reopen the paper-faithful MMT branch;
- use measured target variables as inference inputs;
- tune an anchor on validation or test curves;
- train a neural residual;
- promote a PLC implementation;
- claim trustworthy extrapolation from numerical finiteness alone.

Analytical Formulation

For each operating condition, the forward curve is projected into an explicit offset and ordered sine/cosine coefficients:

$TE(\theta) = \text{offset} + \sum \text{over } k \text{ of } \sin_coefficient_k \text{ times } \sin(k \theta) + \cos_coefficient_k \text{ times } \cos(k \theta).$

The canonical local-order set is

`[1, 3, 39, 40, 78, 81, 156, 162, 240]` .

Each of the nineteen coefficient outputs consists of:

- one offset;
- nine sine coefficients;
- nine cosine coefficients.

Each coefficient is modeled as a complete-quadratic surface of the three operating variables. The design includes:

- a constant term;
- three linear terms;
- three squared terms;
- three pairwise interaction terms.

The operating inputs are standardized using training-only statistics before the quadratic basis is formed. This structure is not a first-principles contact equation. It is an interpretable grey-box analytical representation:

- Fourier orders encode repeatable angular structure;
- coefficient surfaces encode smooth operating-condition dependence;
- the future neural residual can model what this bounded analytical component does not explain.

Reproduction Of Phase 1


The Stage 3 runner independently reloads the frozen curves, reconstructs the training-only coefficient targets, and refits the complete-quadratic surface.

The following comparisons all have zero numerical difference at tolerance 1e-12 :

Reproduced quantity	Result
feature means	exact
feature scales	exact
coefficient matrix	exact
design condition number	exact
test raw MAE	exact
test RMSE	exact
test centered MAE	exact
test centered RMSE	exact
test offset absolute error	exact
test peak-to-peak absolute error	exact
test derivative MAE	exact
retained-amplitude MAE	exact
retained-phase MAE	exact

This proves that the anchor used in Stage 3 is not an approximate recreation or a differently preprocessed variant. It is the Phase 1 PF-A formulation refitted from the frozen training evidence.

Forward Variant Comparison



Forward analytical variant comparison

Six formulations were evaluated on the same 97 forward test conditions.

Rank	Model	Orders	Test MAE [deg]	Role
1	PF-A local	9	0.001807	qualified anchor

Rank	Model	Orders	Test MAE [deg]	Role
2	PF-E reduced	7	0.001823	comparator only
3	PF-G PLC-safe	4	0.001945	comparator only
4	PF-H ONNX-order	3	0.002011	comparator only
5	PF-A paper-order	20	0.002126	comparator only
6	PF-B recovered ONNX	3	0.003047	comparator only

The table uses compact display labels. The complete machine identifiers remain in `stage3_forward_variant_comparison.csv` .

Interpretation

The result rejects a simplistic rule that more harmonics must be better. The twenty-order paper-derived set is worse than the nine-order local set on this frozen forward surface.

The seven-order reduced form is close to PF-A in raw error, but closeness is not enough to replace the anchor. Its order removal changes ripple and harmonic content, and no deployment benefit has yet been demonstrated on the complete multi-index gate.

The four-order PLC-safe core is useful as a low-complexity deployment comparator. It is not promoted because the Stage 3 objective is anchor qualification, not minimum-operation deployment selection.

The recovered ONNX path remains valuable as historical parity evidence but is materially worse as the analytical anchor.

Bootstrap Stability

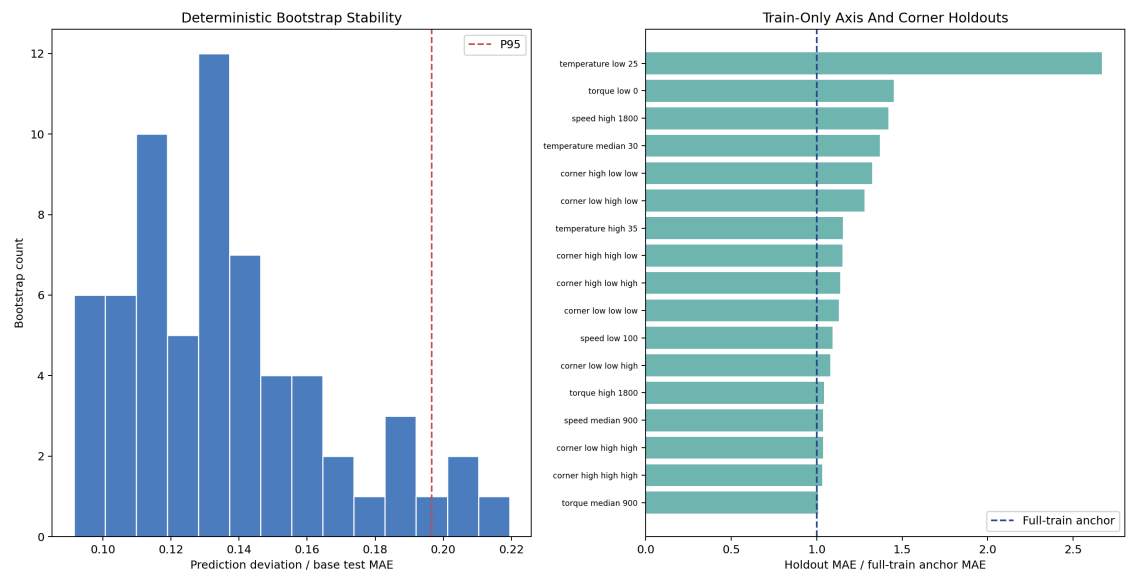
The canonical surface was refitted sixty-four times using deterministic training-only bootstrap samples with seed `314159` .

Diagnostic	Median	P95	Maximum
design condition number	5.391523	5.612715	5.705000
relative coefficient delta	0.013685	0.029912	0.036747
prediction deviation / base MAE	0.131243	0.196634	below gate

The design remains well conditioned across every bootstrap. Coefficients move slightly as expected under resampling, but the induced prediction deviation is small relative to the original test error.

This result supports use of PF-A as a repeatable analytical component. It does not establish uncertainty calibration; Stage 11 addresses uncertainty and physics-trust calibration explicitly.

Train-Only Operating-Condition Holdouts



Bootstrap and holdout diagnostics

Seventeen refits were executed

- low, median, and high torque levels;
- low, median, and high speed levels;
- low, median, and high temperature levels;
- eight geometric operating-space corner populations.

Every holdout fit excludes its evaluation population before estimating feature statistics or coefficients.

Most Informative Failures


Holdout	Support type	Held-out MAE [deg]	Ratio to full anchor
temperature low, 25 C	axis-edge extrapolation	0.004742	2.670
low-speed, high-torque, high-temperature corner	sparse corner	0.004619	about 1.0
speed high, 1800 rpm	axis-edge extrapolation	0.003879	about 1.4
torque low, 0 Nm	axis-edge extrapolation	0.003860	about 1.5
temperature median, 30 C	interpolation	0.002517	about 1.3

The lowest-temperature holdout is the dominant weakness. Temperature is not a minor cosmetic input: removing its lower regime substantially degrades the surface.

The median-temperature holdout raises the condition number to 10.347985 , above the full-fit value but still finite. This indicates that the discrete temperature support is structurally important to the quadratic basis.

The conclusion is not that PF-A is unstable. The conclusion is narrower: PF-A is stable when fitted on the full frozen support, while extrapolating across omitted operating regimes is materially less trustworthy.

Anchor-Corruption Tests



Analytical-anchor corruption sensitivity

Thirty-eight deterministic corruptions cover four families

- coefficient scale;
- phase shift;
- order omission;
- operating-input shift.

Most Sensitive Corruptions

Corruption	Family	Mean deviation from PF-A [deg]
omit low-order group	order omission	0.010982
omit order 1	order omission	0.010928
torque shift, +0.50 training standard deviations	input shift	0.010322
torque shift, -0.50 training standard deviations	input shift	0.010090
torque shift, +0.25 training standard deviations	input shift	0.005132
torque shift, -0.25 training standard deviations	input shift	0.005074
scale all coefficients by 1.10	coefficient scale	0.005024
scale all coefficients by 0.90	coefficient scale	0.005024
phase shift, +15 deg	phase	0.002864
phase shift, -15 deg	phase	0.002843

Physical And Modeling Meaning

The lowest orders carry the dominant slow angular structure. A residual model must not be allowed to erase or arbitrarily rewrite this content without an explicit diagnostic.

Torque-input corruption is nearly as damaging as low-order omission. This supports two later requirements:

- 1. preserve input provenance and scaling exactly;
- 2. monitor analytical versus neural contribution by harmonic band.

Phase corruption is detectable but less damaging than low-order removal at the tested magnitudes. This does not make phase unimportant. It means the tested PF-A error surface is more sensitive to missing dominant content and operating-input displacement than to a uniform modest phase rotation.

Deployable Validity Envelope

The envelope is derived only from the 675 training conditions.

Training Axis Bounds

Input	Minimum	Mean	Maximum	Scale
signed torque [Nm]	-1800.610	-914.065	0.114	547.043
absolute speed [rpm]	100.000	934.519	1800.001	534.978
temperature [C]	24.259	30.850	37.642	4.006

The density threshold is the training-only P95 leave-one-out nearest distance in standardized operating space: 0.303207 .

Tiers

Tier	Definition	Runtime decision
Core	inside all training bounds and within density threshold	PF-A may be used as qualified anchor
Sparse or corner	inside bounds but farther than density threshold	low-trust anchor; monitor or route conservatively
Extrapolation	outside one or more training bounds	fallback or explicit review

The complete machine labels are supported_core , supported_sparse_or_corner , and unsupported_extrapolation .

Observed Population

Split	Core	Sparse or corner	Extrapolation
train	675	0	0
validation	186	7	1
test	90	6	1

The twenty-seven-point minimum/mean/maximum envelope grid is numerically finite. All 966 eligible forward predictions are also finite.

Numerical finiteness is necessary but not sufficient. The deployment rule is therefore based on support classification, not only on a finite-number check.

Exit-Gate Results

Gate	Result
frozen split signature matches	pass
refit uses training partition only	pass
Phase 1 reproduction	pass
explicit offset and sine/cosine coefficients	pass
condition-number stability	pass
bootstrap coefficient stability	pass
bootstrap prediction stability	pass
required forward variant roster	pass
train-only axis and corner holdouts	pass
all corruption families	pass
deployable validity envelope	pass
finite predictions on every eligible forward condition	pass

Result

12 / 12 gates pass.

What Stage 3 Proves

- Stage 3 proves that
- the repository can refit PF-A without hidden validation or test fitting;
 - the refit is exactly reproducible against Phase 1;
 - the quadratic operating-condition basis is numerically stable on the full training population;
 - coefficient and prediction variation under bootstrap resampling is bounded;
 - the local nine-order formulation is the strongest tested analytical anchor;
 - the anchor remains finite over the complete eligible forward population;
 - support-aware deployment tiers can be computed causally from training operating inputs;
 - low-order content and torque input integrity are critical.

What Stage 3 Does Not Prove

Stage 3 does not prove that

- PF-A is the best final predictor;
- PF-A is physically exact;
- a residual neural network will improve it;
- every point inside axis bounds is equally trustworthy;
- numerical output outside the support envelope is reliable;
- the reduced or PLC-safe variants are deployment-ready;
- the future residual should modify every harmonic band;
- a physics-guided model beats a parameter-matched data-only model.

Stage 4 Design Consequences

Stage 4 is the Data-Only Residual Capacity Ladder. It must use PF-A as a frozen, bounded anchor and answer a deliberately non-physics question first:

How much of the remaining error can an ordinary residual network learn before any physics-guided loss is introduced?

The Stage 4 campaign should therefore

1. keep PF-A frozen in the primary residual arms;
2. expose analytical and residual contributions separately;
3. include zero-residual and direct-black-box controls;
4. test a small capacity ladder under matched budgets;
5. report residual energy and harmonic-band projection;
6. preserve the support tiers in evaluation;
7. fail closed on unsupported extrapolation;
8. avoid crediting capacity or normalization gains to physics.

Only after this data-only capacity floor is known can Stage 5 determine whether complex harmonic coefficient supervision adds real value.

Durable Artifacts

Implementation

- Stage 3 analysis runner;
- Stage 3 independent validator.

Machine Evidence

Under `output/analysis/wave_5_2r/stage3_analytical_anchor_reproduction_and_stress_tests/`

- refitted PF-A surface;
- explicit coefficient-surface table;
- Phase 1 reproduction comparison;
- six-variant comparison;
- bootstrap repeat and target diagnostics;
- holdout diagnostics;
- corruption diagnostics;
- per-condition validity-envelope assignments;
- validity-envelope summary;
- twelve-gate exit summary.

Human Evidence

- three report plots;
- script usage note;
- project usage-guide entry;
- this Markdown report;
- validated PDF companion.

Conclusion

PF-A has earned a precise role.

It is not a full physical solution, and it is not automatically safe outside measured support. It is a stable, reproducible, inspectable analytical component that captures the dominant forward harmonic structure and smooth operating-condition dependence better than the tested analytical alternatives.

That makes it a sound anchor for residual learning. The lowest-temperature holdout, low-order omission, and torque-input corruption results also show exactly where the future hybrid must remain cautious.

Stage 4 is authorized to measure residual-network capacity on top of this qualified anchor. No physics-guided advantage may be claimed until that data-only residual baseline exists.